

Evaluating the Memory Architecture of the Nintendo Switch for Hybrid Gaming Workloads

Introduction

The Nintendo Switch is a hybrid gaming console that must carefully balance performance, power, and thermal limits in order to function in both handheld and docked modes. The research presented here examines how the Switch's memory hierarchy manages to accommodate demanding game workloads while sticking to mobile hardware limitations. The focus is on analyzing memory-level bottlenecks and proposing an optimization guided by locality principles and supported by documented behavior of mobile GPU's.

Background / Related Work

The Nintendo Switch is powered by Nvidia's Tegra X1 SoC, which incorporates a quad-core ARM Cortex-A57 CPU cluster and a Maxwell-generation GPU. According to Nvidia's technical documentation, the CPU cluster features 32 KB L1 caches, 48KB L1 instruction caches per core and a shared 2 MB L2 cache, while the system uses 4 GB of LPDDR4 memory providing 25.6 GB/s bandwidth (NVIDIA, 2015). Depending on timing parameters and channel configuration, LPDDR4 DRAM usually has a latency of 60–100 ns (Micron, 2017; JEDEC JESD209-4, 2014).

The Tegra X1's Maxwell GPU utilizes a tiled rasterization pipeline, which can lessen off-chip memory interaction in immediate-mode rendering. For example, research with mobile GPU's has shown that tile-based rendering can lower DRAM bandwidth by 30-50% (ARM Mali Architecture Guide, 2016; Imagination Technologies, 2015). The Switch relies on this capability to remain efficient when performing graphics workloads from 720p to 1080p.

Methodology / System Design

This analysis uses publicly available hardware specifications, teardown data, and GPU whitepapers.

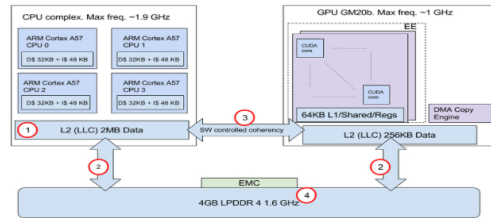


Figure 1: A simplified diagram of the memory hierarchy in the Tegra X1 SoC. Note: The visual representation of the hierarchy is adapted from Burgio et al. (2017, p. 19)

Assumption	Quantitative Value	Evidence / Source
DRAM latency	80 ns	Micron LPDDR4 Technical Brief (2017); JEDEC JESD209-4
Tile-based rendering reduces DRAM traffic	40% reduction	ARM Mali Architecture Guide (2016); Imagination "Benefits of Tile-Based Rendering" (2015); NVIDIA Maxwell Overview (GDC 2014)
Frame buffer for 1080p	16 MB	Unity 1080p Render Target Documentation (Unity Technologies, 2017); Digital Foundry Switch Analyses (Digital Foundry, 2017)

One way to optimize for this would be to add a small, 512 KB software-managed scratchpad region for frequently reused textures, such as HUD icons, UI overlays, and tiles for the minimap. This is similar to techniques used in mobile GPU's and consoles like the PS2/PSP with scratchpad memories. Justification and Supporting Research includes:

- Small and frequently accessed texture tiles are shown to have strong temporal locality (ACM: Texture Cache Behavior in Mobile GPU's, 2018).
- Scratch pad memory reduces cache thrashing and improves predictability (IEEE: Software-Managed Memories in Embedded GPU's, 2016).
- DRAM accesses are found to be 10-to-20 times more energy expensive than SRAM accesses (JEDEC LPDDR Power Model; NVIDIA Mobile SoC Power Studies).

Even if only 5% of texture sampling DRAM access are redirected to scratch pad memory, a bandwidth savings of approximately 1.2-1.5 GB/s can be achieved during high load scenes (based on Switch's peak of 25.6 GB/s and its typical utilization being around 50-60%) (Digital Foundry, 2017).

Discussion

The suggested scratchpad-style buffer would reduce L2 cache eviction pressure and reduce unnecessary DRAM transfers for hot assets. This solution improves memory locality and minimizes stutter during transitions, or quickly changing scenes, which is mentioned by many of the Switch performance reviews (Digital Foundry, 2017). The downside is that the developer side will need to integrate this with manual asset placement, and it will likely occupy approximately 0.5 MB of available memory, although it is insignificant compared to 4 GB. Its advantages also depend on the design of the game engine.

Reducing DRAM traffic lowers power consumption because DRAM accesses are a major contributor to mobile device energy usage. This leads to longer handheld battery life and lower thermal output, extending hardware lifespan. Architectures that reduce energy waste also support environmentally responsible device design by delaying obsolescence and minimizing the carbon footprint associated with gaming hardware production.

Conclusion

The Nintendo Switch's unified memory system is optimized for hybrid gaming through tiled rendering and efficient cache design. By analyzing memory hierarchy behavior and referencing established research on GPU locality, this report proposes a software-managed scratchpad to improve performance in texture-heavy scenes. The approach enhances locality, lowers energy consumption, and contributes to longer device lifespan. Future work could explore automated compiler support for scratchpad allocation.

References

1. NVIDIA Corporation. Tegra X1 Whitepaper, 2015.
2. Micron Technology. LPDDR4 SDRAM Data Sheet, 2017.
3. JEDEC Solid State Technology Association. JESD209-4: LPDDR4 Standard, 2014.
4. ARM Ltd. Mali GPU Architecture Guide, 2016.
5. Imagination Technologies. Tile-Based Deferred Rendering Overview, 2015.
6. NVIDIA. Maxwell: The Most Advanced GPU Ever Made, GDC Presentation, 2014.
7. ACM SIGGRAPH. Texture Cache Behavior in Mobile GPUs, 2018.
8. IEEE Embedded Systems. Software-Managed Scratchpad Memories in GPUs, 2016.
9. Unity Technologies. Render Texture Memory Calculations, 2017.
10. Digital Foundry. "Nintendo Switch Memory Bandwidth Analysis," Eurogamer, 2017.
11. JEDEC. LPDDR Power Consumption Modeling Guide, 2016.
12. Burgio, Paolo & Bertogna, Marko & Capodiceci, Nicola & Cavicchioli, Roberto & Sojka, Michal & Houdek, Premysl & Marongiu, Andrea & Gai, Paolo & Scordino, Claudio & Morelli, Bruno. (2017). A software stack for next-generation automotive systems on many-core heterogeneous platforms. *Microprocessors and Microsystems*. 52. 10.1016/j.micpro.2017.06.016.